

# Senior Machine Learning Engineer — Training Infrastructure (US)

Company: Amazon

Role family: MLE | Seniority: Senior | Location: United States (role location varies by team)

Source (public posting): <https://www.amazon.jobs/en/jobs/2668768/sr-machine-learning-engineer>

Cluster demand (ground-truth): {"MLE": 0.7, "SWE": 0.25, "DS": 0.05, "QD": 0.0, "QR": 0.0}

Reconstruction method: normalized sections + paraphrased responsibilities/qualifications for QA use.

## ROLE SUMMARY

You will lead engineering work that enables large-scale model training and experimentation. The emphasis is on building stable, efficient training systems and collaborating with scientists/engineers to ship production-ready ML capabilities.

## KEY RESPONSIBILITIES

- Build and evolve distributed training infrastructure that scales to large models and GPU/accelerator hardware.
- Develop tooling for data ingestion, feature/label pipelines, and repeatable training runs.
- Improve reliability and performance (throughput, cost, latency) of training and evaluation workflows.
- Establish best practices for reproducibility, model/version tracking, and deployment handoffs.
- Partner across ML, product, and platform teams to define requirements and deliver systems that unblock research-to-production.

## REQUIRED QUALIFICATIONS

- 5+ years professional software engineering experience (Python/Java/C++ or similar).
- 3+ years building ML systems or ML infrastructure in production.
- Strong knowledge of distributed systems, debugging, and performance profiling.
- Experience with GPUs/accelerators, containerized compute, and CI/CD for ML.

## PREFERRED QUALIFICATIONS

- Experience with large model training and evaluation at scale.
- MLOps patterns: monitoring, drift detection, retraining orchestration, and model governance.